

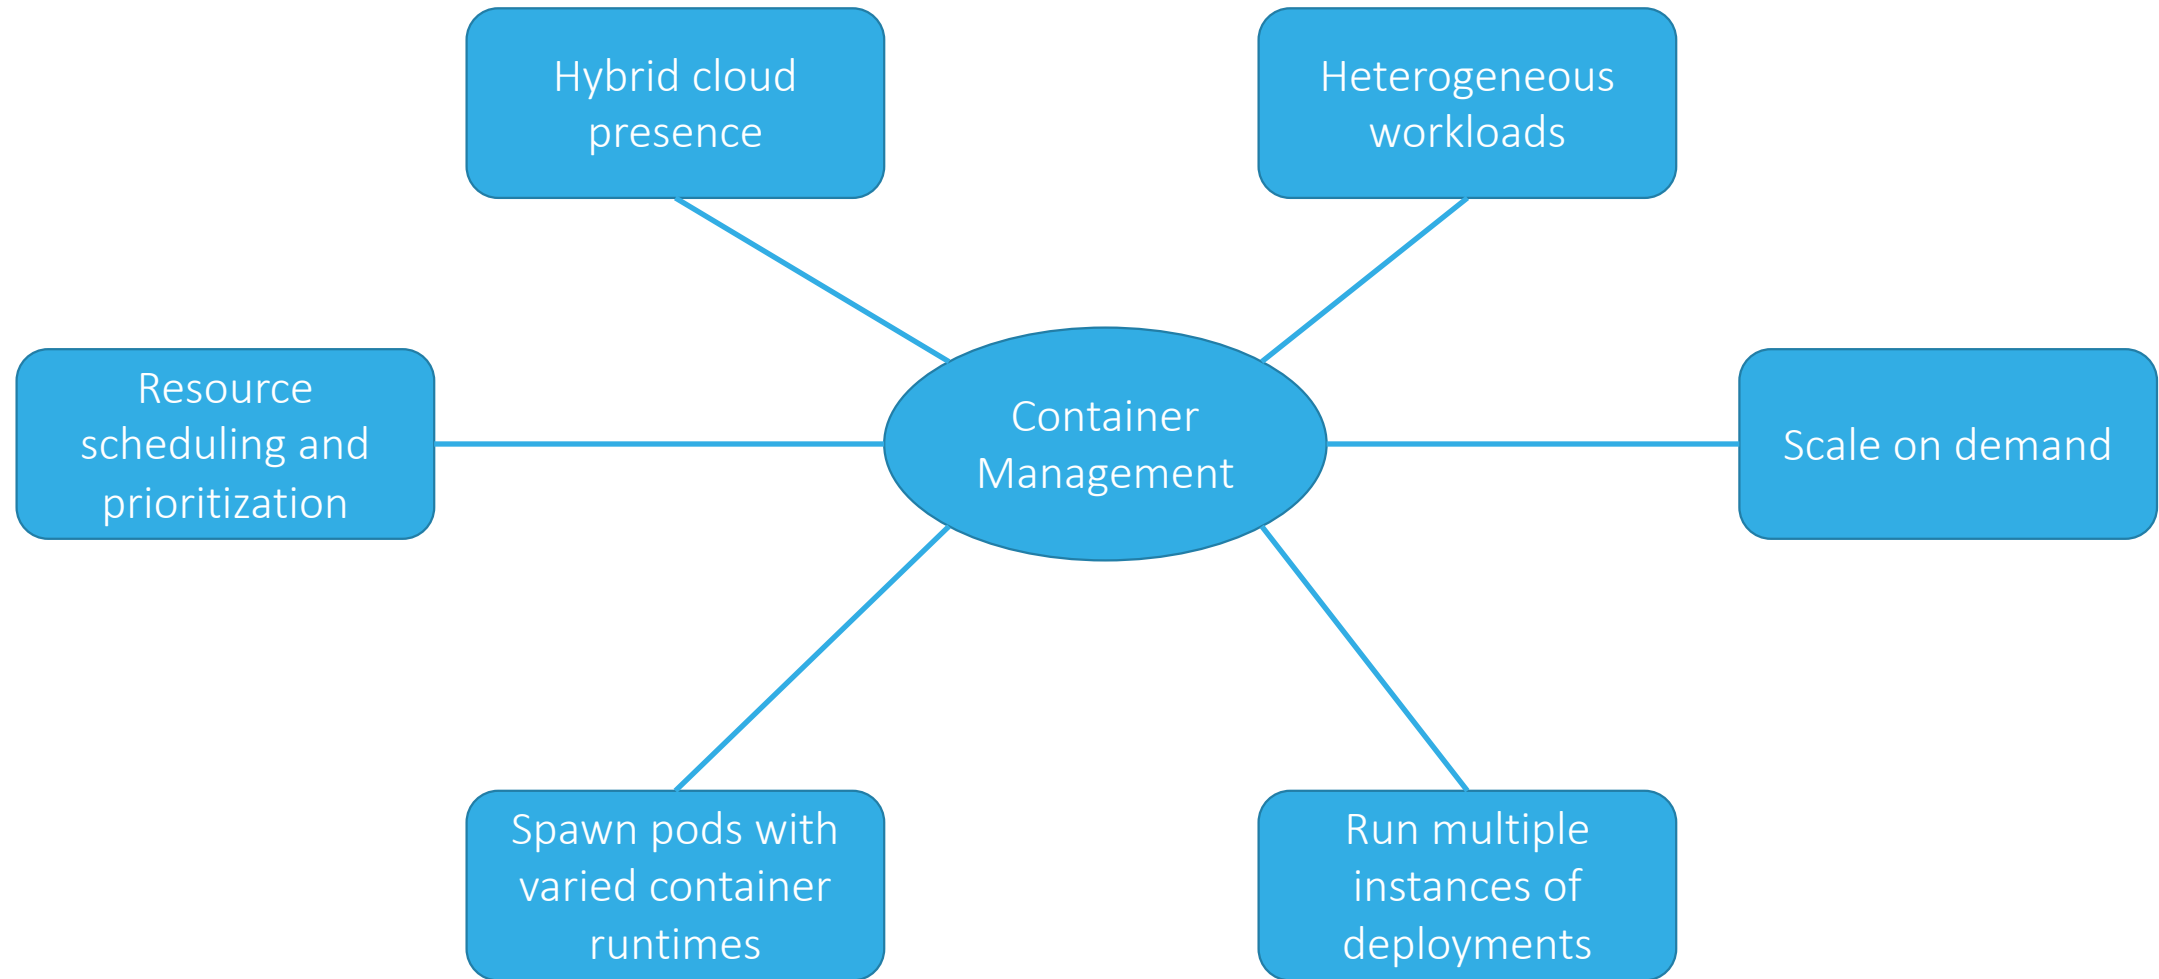
Optimizations and design considerations for Kubernetes cluster for Large Scale

RaviShankar KS
Aug, 2019

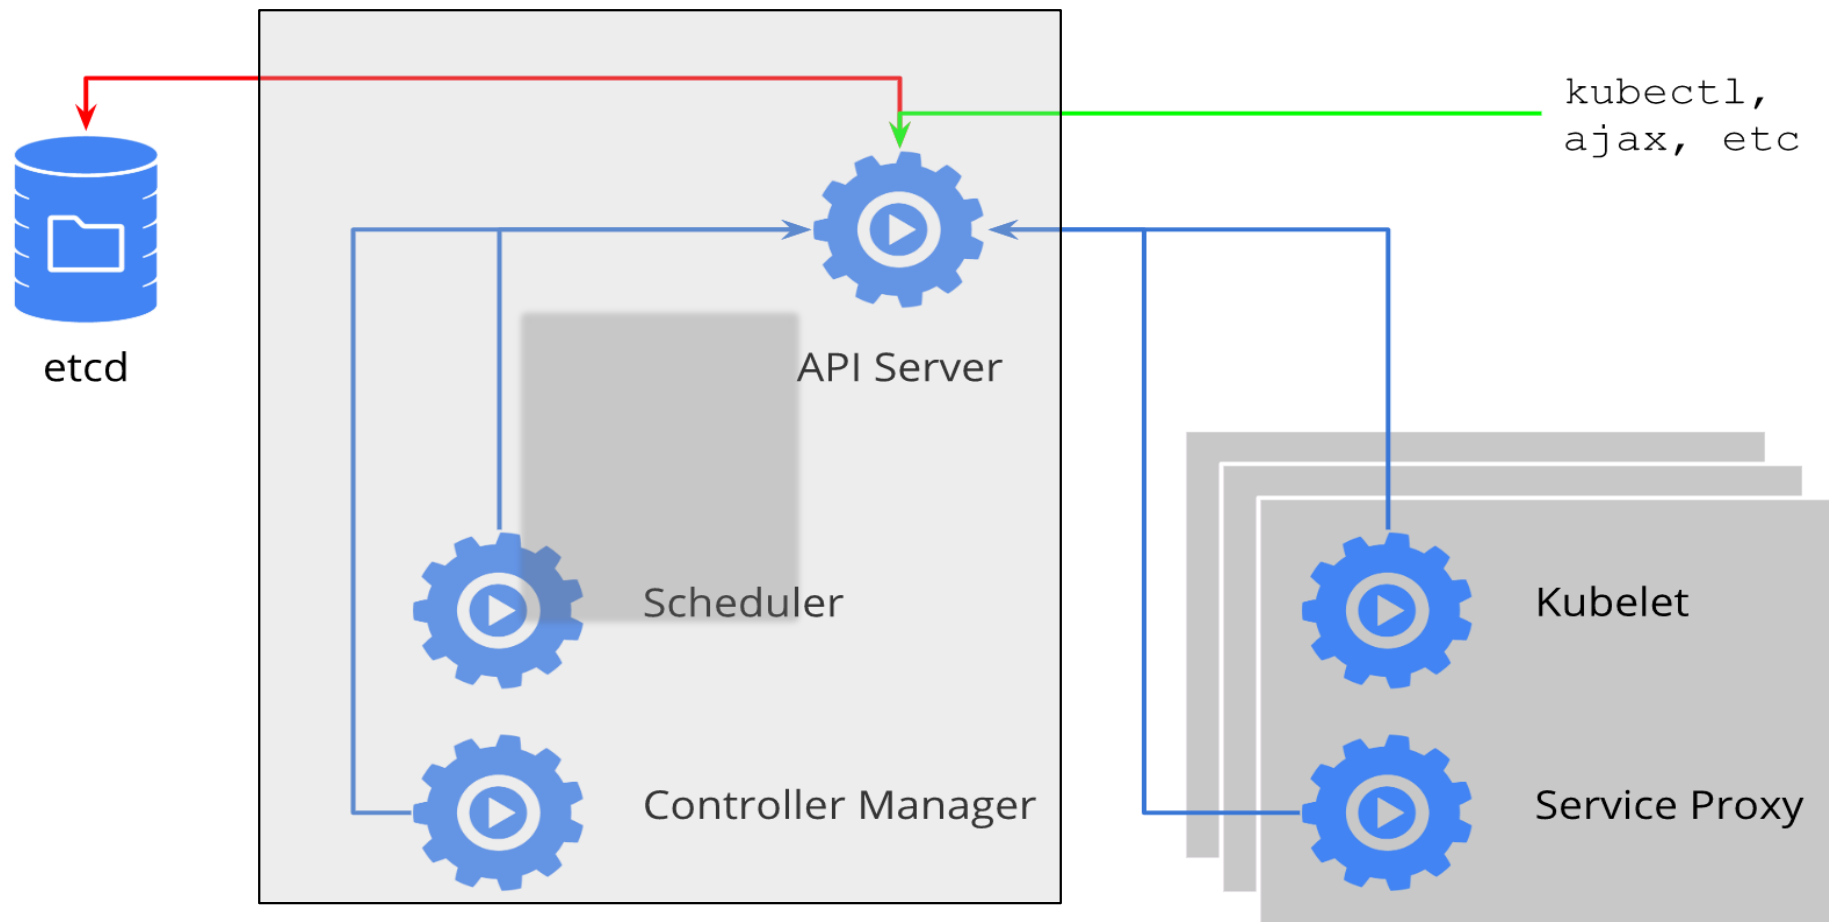
Outline

- Expectations
- Deep Dive
 - K8s components overview
 - SLAs and tuning
 - Dynamic Autoscaling
 - Monitoring
- Few more Lessons Learnt trivia of interesting use-cases and solution
- Q&A

Expectations



Kubernetes Architecture



Kubernetes system components

etcd

- Persistent Key value store
- Needs Quorum/Leader to operate

apiserver

Security

- Admission Control – RBAC
- Authentication

Objects

- Caching
- Validation
- CRUD operations

Docker containers

- Logging
- Exec

Controller manager

- Watch/List k8s resources
- Uses Apiserver
- Moves resources from current state to desired state

Scheduler

- Watch Pending pods
- Eliminate nodes which are not fit for request to be fulfilled
- Rank Nodes

Kubelet + kube-proxy

- Node Level admission control
- Communicate with container through CRI
- Node metrics
- Report readiness / liveness

kube-proxy

- Update iptables with service endpoints

opportunities for optimization

- Keep K8s production grade and always running
 - Parameter tuning and trade-offs
 - Process panic
 - Node issues

Parameter tuning

apiserver	
# of k8s master nodes	3-4 Multi-zonal
Rate limit	--max-requests-inflight [400]
Choosing the correct Storage backend	[etcd3]
Feature gates	cpumanager debugcontainer
Caching	--target-ram-mb --event-ttl --audit-log-maxage
Auditing	--auditing*
# requests to etcd	--etcd-count-metric-poll-period duration
etcd	--etcd-compaction-interval
Expose Kubernetes service	--kubernetes-service-node-port

Parameter tuning

Node Level	
<ul style="list-style-type: none">• Docker daemon hangs• Kernel crashes• Kubelet crash	Cron jobs
Monitoring	read-only-port [10250]
Docker fs issues	use overlay2
Port conflicts	net.ipv4.ip_local_port_range = [32768 61000] net.ipv4.ip_local_reserved_ports = [30000 – 32000]

etcd	
Quorum size	3,5,7 Multi-zonal
etcd sync	ETCD_ELECTION_TIMEOUT ETCD_HEARTBEAT_INTERVAL

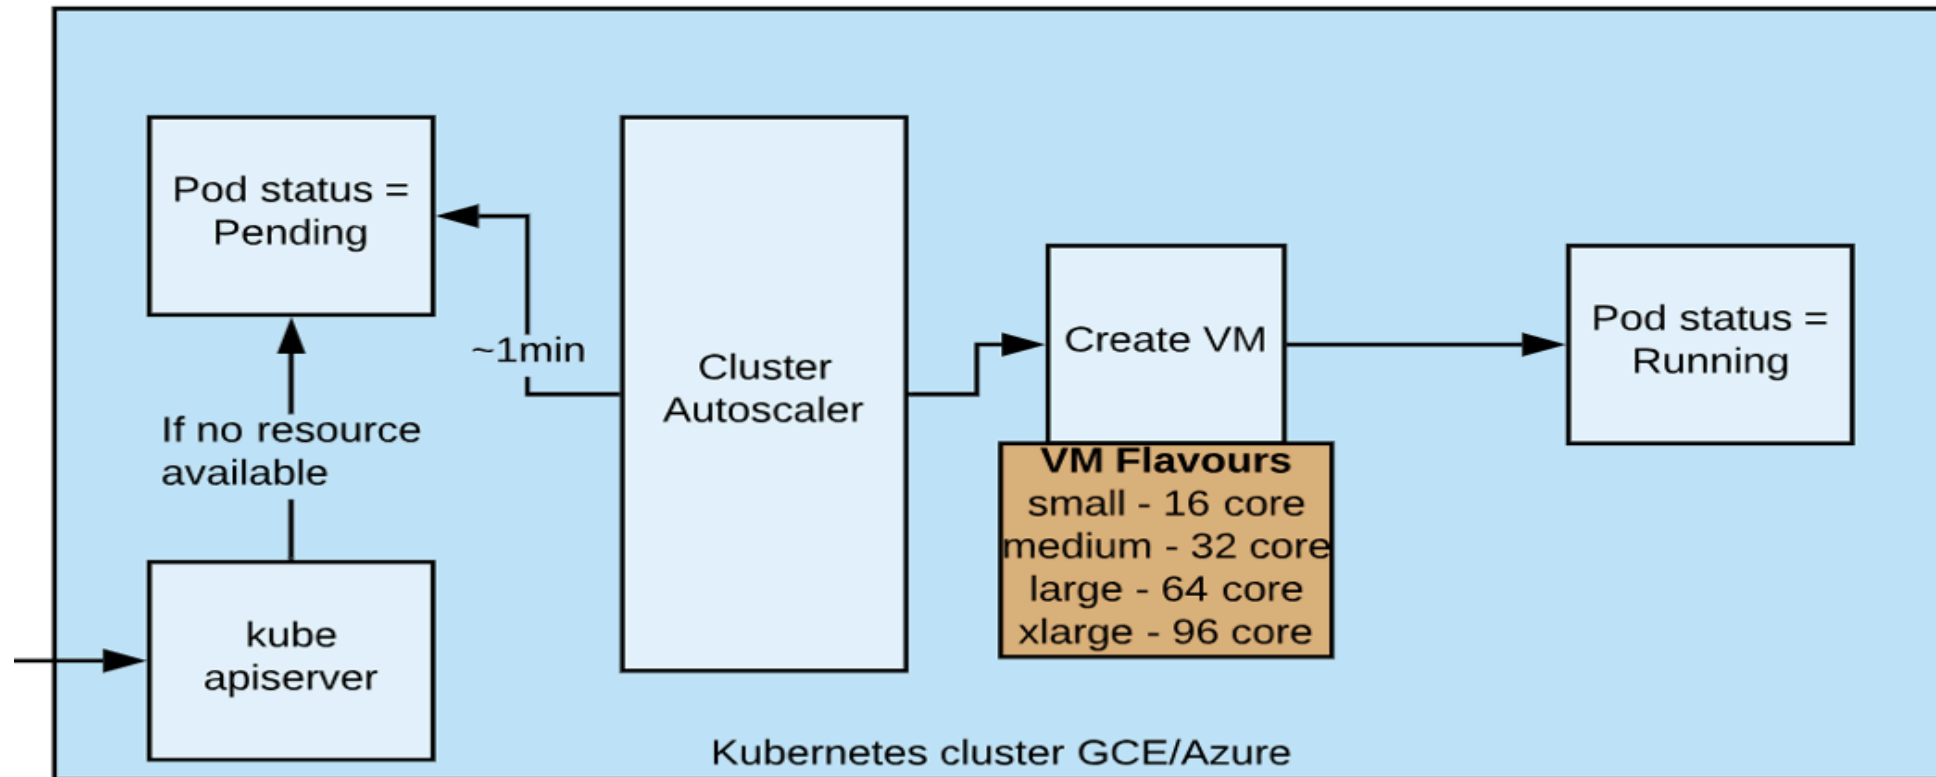
general	
Networking choices	Calico / Flannel

controller manager	
Node monitoring	--node-monitor-period --node-monitor-grace-period
Parallelism – number of workers	--concurrent*
Caching	*cache-size*
Rate limit	--kube-api-qps --kube-api-burst
GC	--terminated-pod-gc-threshold

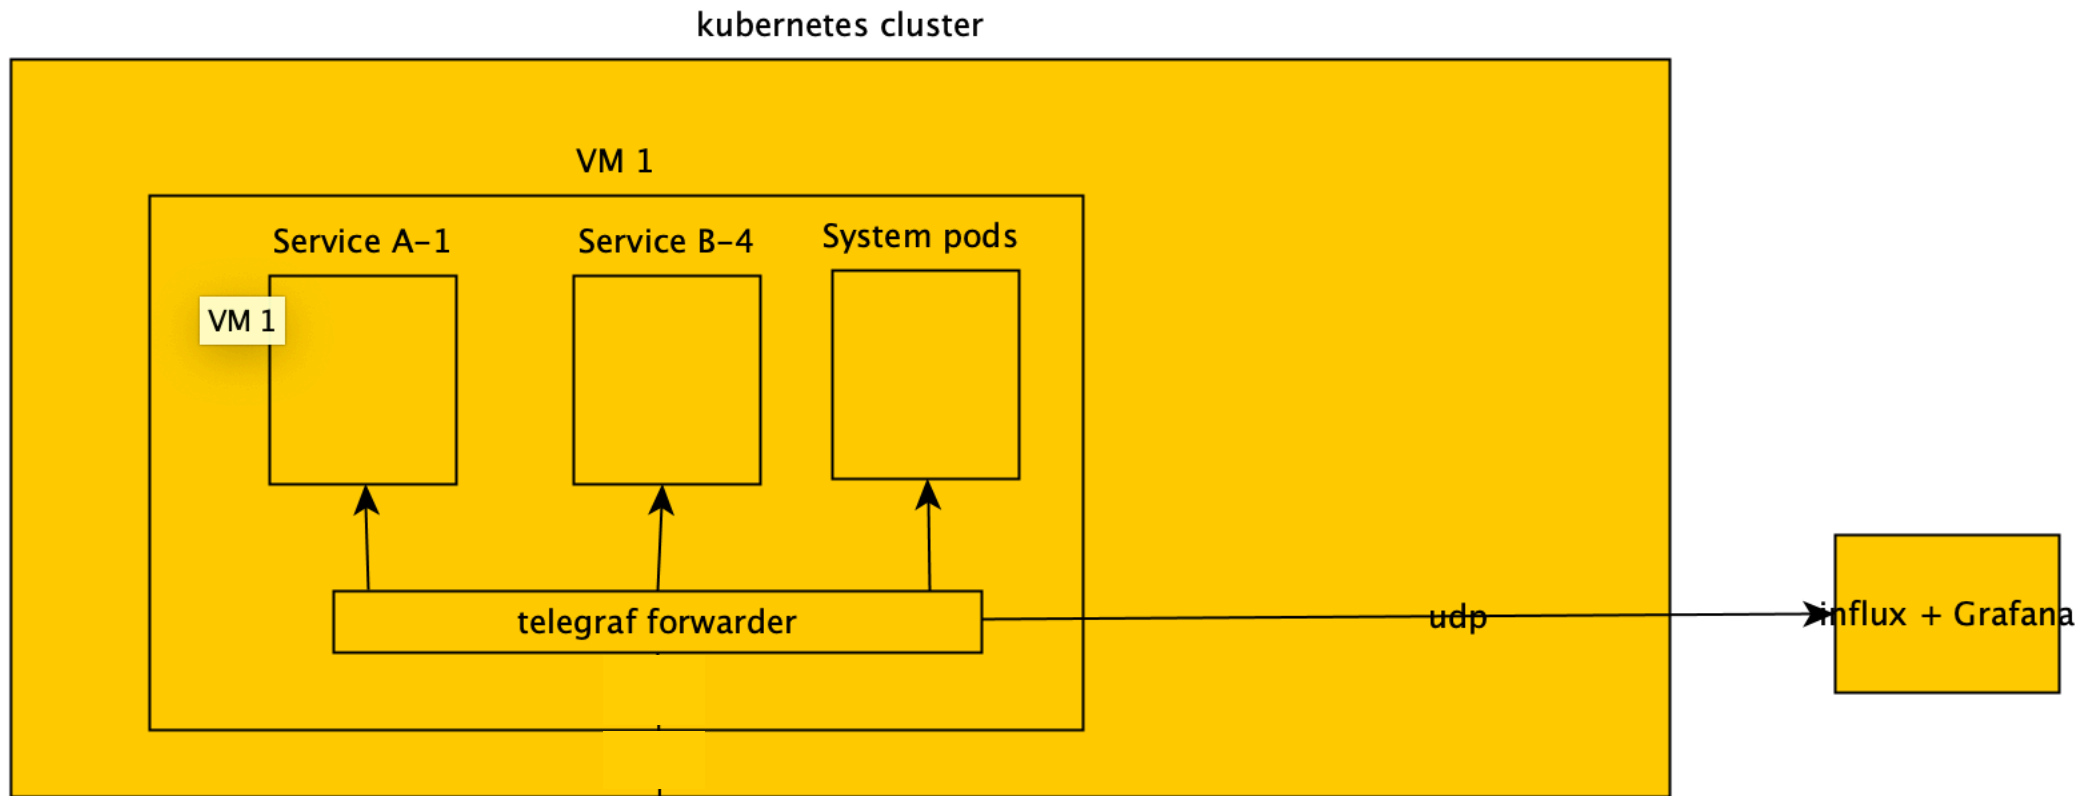
opportunities for optimization

- Capacity Allocation
 - granularize minions into flavors to avoid node wastage, reduce cost
- Auto-scaling for public cloud (Azure / GCP)
 - Modify kubernetes cluster-autoscaler for node autoscaling
 - Observe Pitfalls of API usage and rate limiting
- Logging / monitoring
 - Every metric matters

Capacity Allocation and autoscaling



Monitoring



Application Monitoring

1. API calls
2. DB conn
3. Health

Node/VM Monitoring

1. Disk Space
2. Disk Mounts


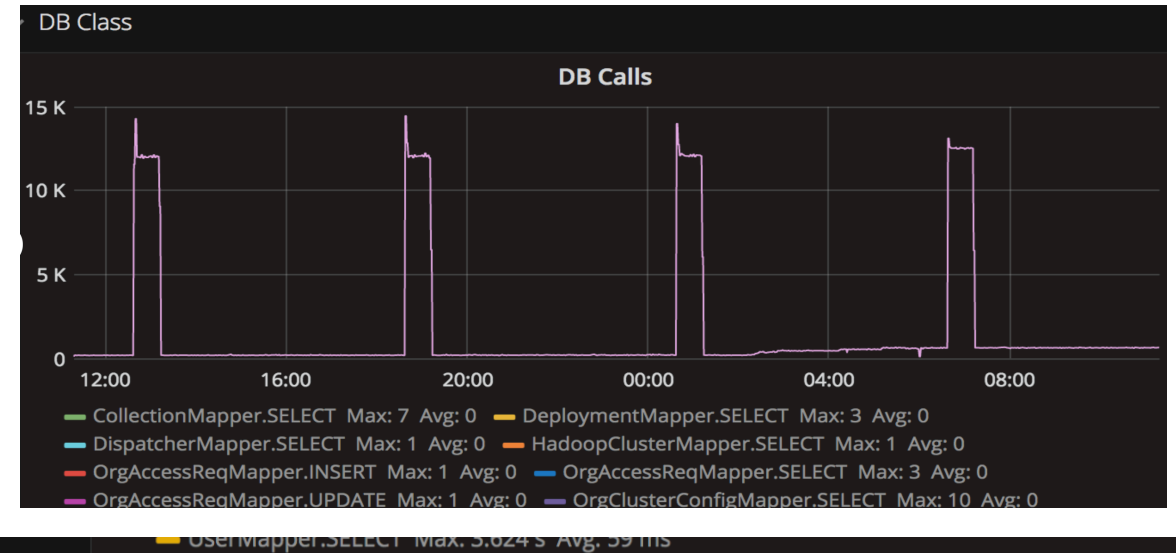

Device Monitoring

1. GPU
2. CPU/Memory
3. Disk IO

Cluster Monitoring

1. ApiServer
2. Controller
3. Scheduler
4. Kubelet
5. Etcd
6. DB Health

Search for patterns



Search for patterns



More lessons learnt....

- Using managed platforms – know limits
 - IP space restrictions over vpc
 - Image prepulls – not supported with coreos (GCE)
 - Limits for # pods, # node-groups not documented
- For persistent external storage
 - Using mount propagation technique
 - Kubernetes Volume Types
- Multi-zonal or multi-regional cluster setup for HA
- IP CIDR Allocation
 - Avoid conflicts
 - Keep room for long term increase in pods/services
- Using ingress endpoints – avoid more number of NodePorts/ClusterIP service types

Q&A

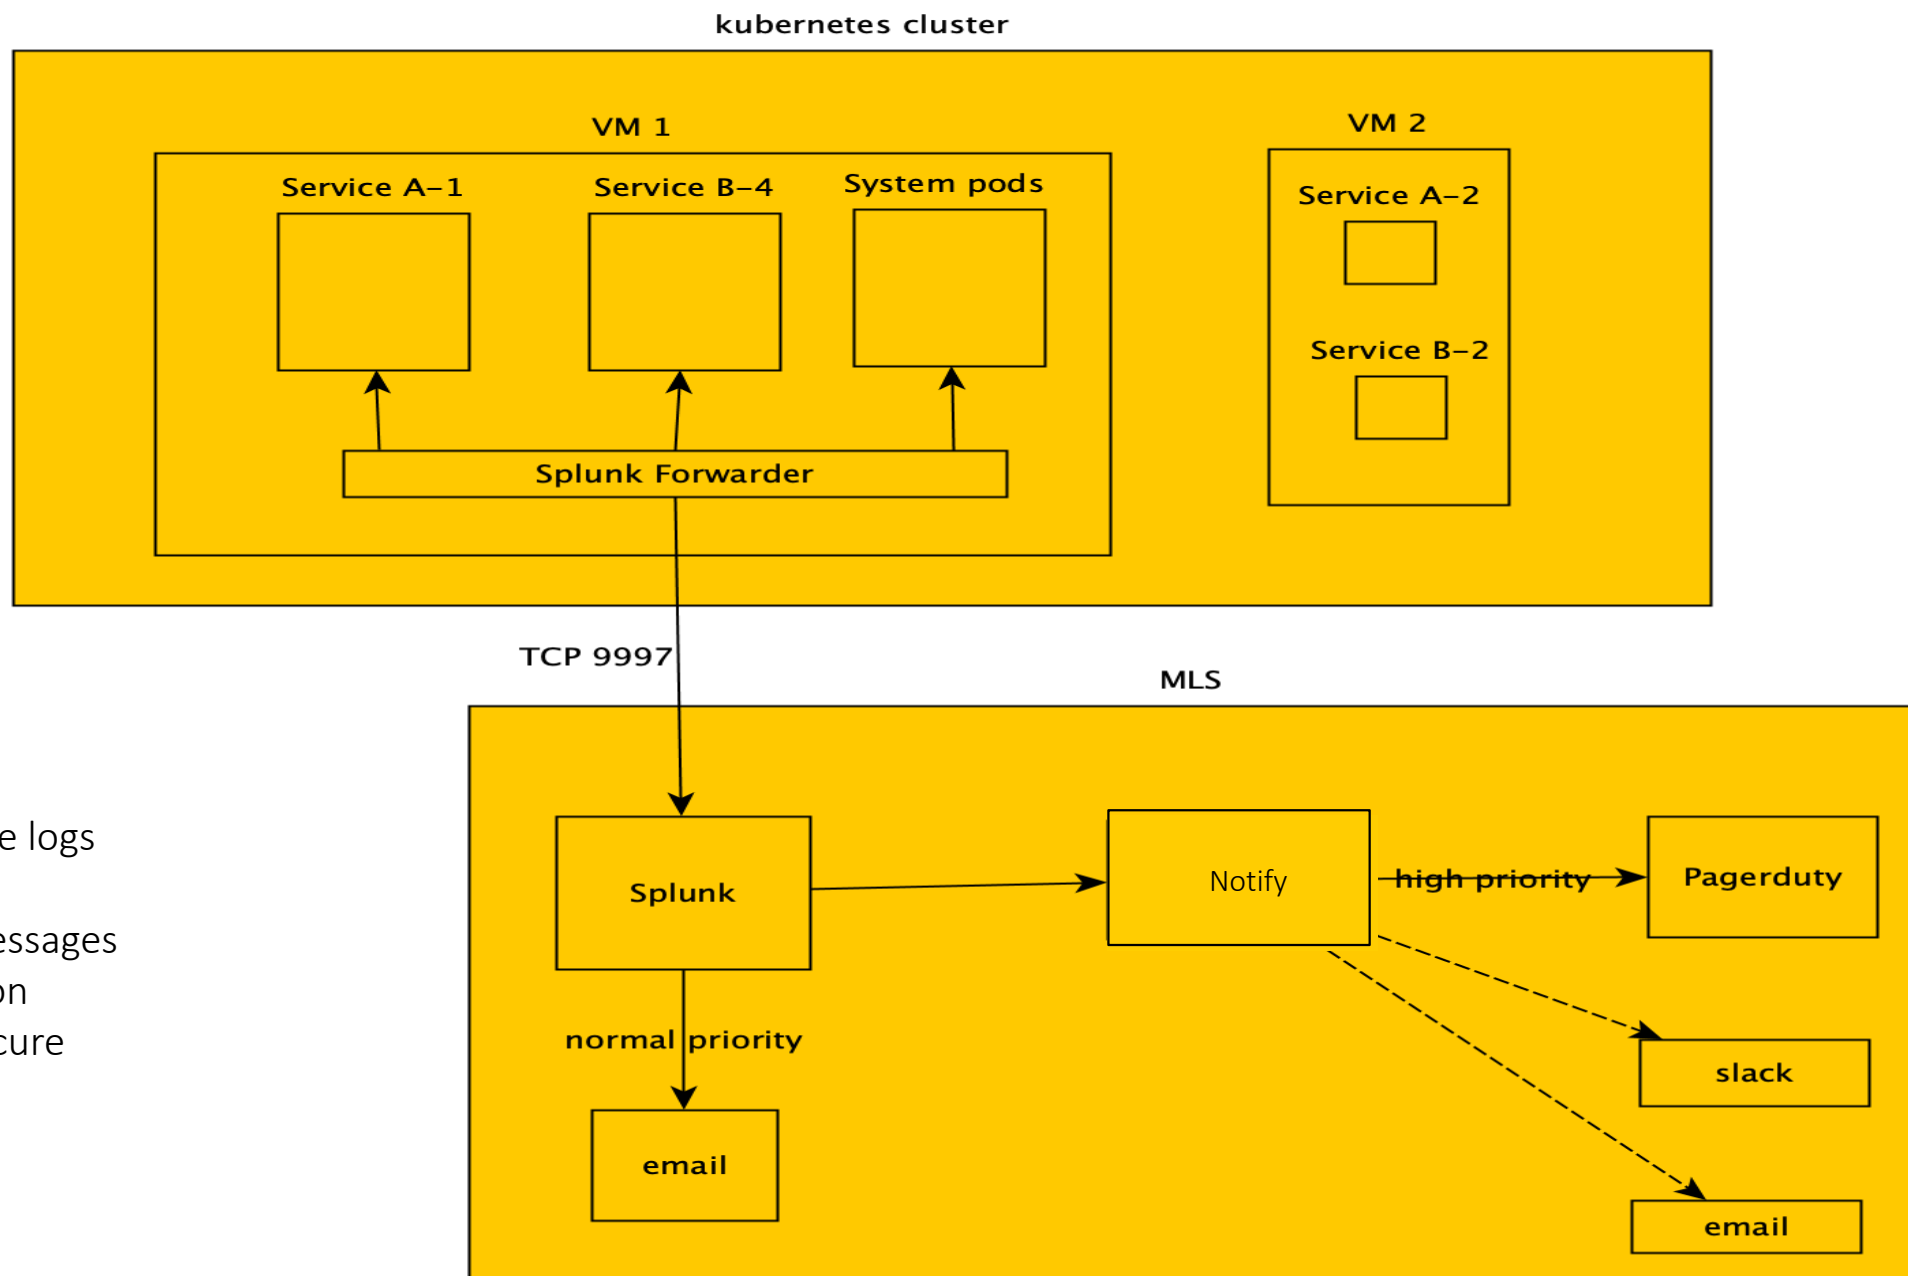
References

- <https://www.slideshare.net/harryzhang735/kubernetes-beyond-a-black-box-part-1>
- <https://www.slideshare.net/harryzhang735/kubernetes-beyond-a-black-box-part-2>
- <https://www.slideshare.net/applatix/webcast-making-kubernetes-production-ready-78130374>
- <https://www.slideshare.net/LCChina/scale-kubernetes-to-support-50000-services>
- <https://unix.stackexchange.com/questions/111858/randomizing-the-source-port-for-new-connections/111869#111869>

Addendum slides

Design considerations / Features	Kubernetes	Docker Swarm
Open Source	Highly modularized with tons of plugins	Less modularized
Stable setup in hybrid clouds	Highly stable, if configured with correct settings for individual components	Limited Fault Tolerance. X
Scale On-demand	Cluster-autoscaler module	Not natively supported X
Resource scheduling and prioritization	Native support	Native support
Heterogeneous workloads	NVIDIA plugins available	Lot of scripting Effort X
Run multiple replicas with HA	Available as deployments	Native support
Access Control	Native RBAC support	Only available in EE X
Other features – privileged mode, mount propagation, NFS support, cpu pinning	Natively supported or plugins available	Complex to setup X
Support base	Large active community	Not much traction , though still pursued. X
Learning curve	Steep learning curve during setup/maintenance phase	Fully Integrated to docker Engine/CLI
Upgrade to higher versions	Not Easy	Easy to upgrade
Post setup maintenance phase	Various choices for logging and monitoring	3 rd party only

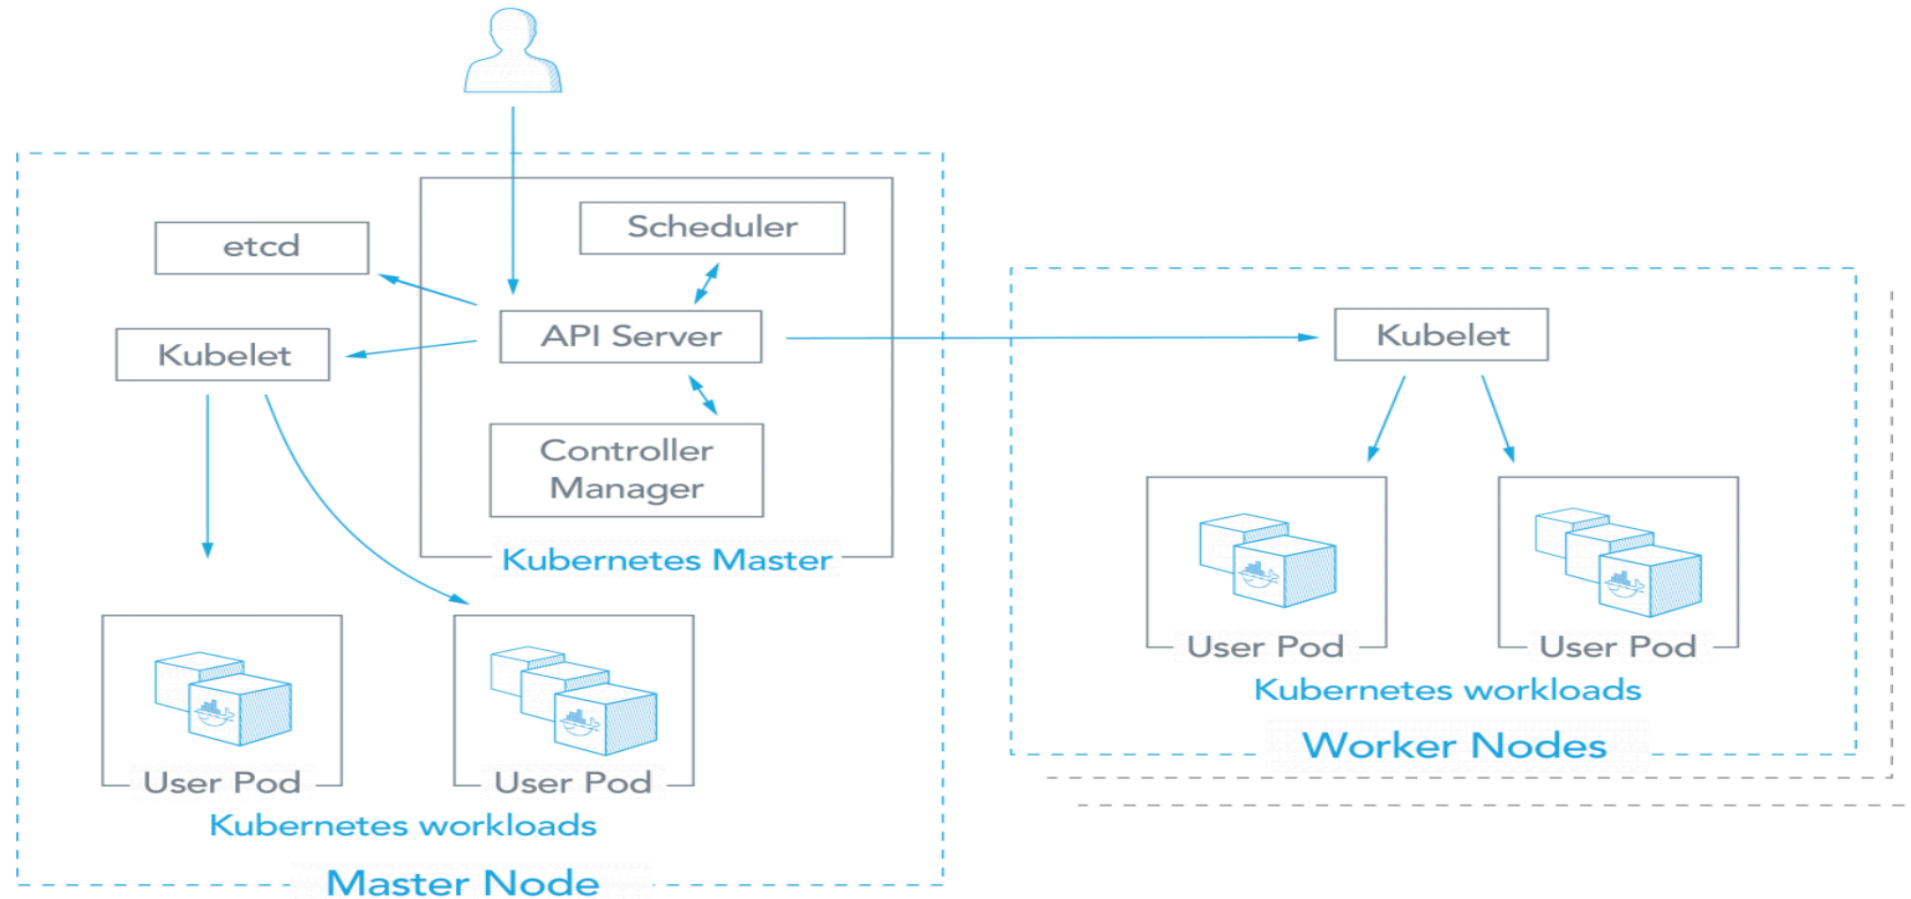
Logging



More Lessons Learnt

- Why not use GKE ??
 - IP space restrictions
 - Image prepulls – not supported with coreos
- GCS mounts on GPU
 - Use mount propagation
- Multi-zonal or multi-regional cluster setup for HA
- IP CIDR Allocation
 - Avoid conflicts
 - Keep room for long term increase in pods/services

Kubernetes Internals



API Server: management hub for Kubernetes
Scheduler: places a workload on the appropriate Node
Controller Manager: scales workloads up/down
etcd: stores configuration data which can be accessed by API Server

Kubelet: Receives pod specifications from API Server, updates Nodes
Master Node: places workloads on Nodes
Worker Nodes: receives requests from Master Nodes and dispatches them
User Pod: a group of containers with shared resources